

D1

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\1_fiat500_VehicleSelection_Dataset.csv")
df
```

Out[2]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1.0	lounge	51.0	882.0	25000.0	1.0	44.907242	8.6115598
1	2.0	pop	51.0	1186.0	32500.0	1.0	45.666359	12.241889
2	3.0	sport	74.0	4658.0	142228.0	1.0	45.503300	11.417
3	4.0	lounge	51.0	2739.0	160000.0	1.0	40.633171	17.634609
4	5.0	pop	73.0	3074.0	106880.0	1.0	41.903221	12.495650
...
1544	NaN	NaN	NaN	NaN	NaN	NaN	NaN	lenç
1545	NaN	NaN	NaN	NaN	NaN	NaN	NaN	conu
1546	NaN	NaN	NaN	NaN	NaN	NaN	NaN	Null valu
1547	NaN	NaN	NaN	NaN	NaN	NaN	NaN	fi
1548	NaN	NaN	NaN	NaN	NaN	NaN	NaN	sear

1549 rows × 11 columns

In [3]: `df.head(10)`

Out[3]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1.0	lounge	51.0	882.0	25000.0	1.0	44.907242	8.611559868
1	2.0	pop	51.0	1186.0	32500.0	1.0	45.666359	12.24188995
2	3.0	sport	74.0	4658.0	142228.0	1.0	45.503300	11.41784
3	4.0	lounge	51.0	2739.0	160000.0	1.0	40.633171	17.63460922
4	5.0	pop	73.0	3074.0	106880.0	1.0	41.903221	12.49565029
5	6.0	pop	74.0	3623.0	70225.0	1.0	45.000702	7.68227005
6	7.0	lounge	51.0	731.0	11600.0	1.0	44.907242	8.611559868
7	8.0	lounge	51.0	1521.0	49076.0	1.0	41.903221	12.49565029
8	9.0	sport	73.0	4049.0	76000.0	1.0	45.548000	11.54946995
9	10.0	sport	51.0	3653.0	89000.0	1.0	45.438301	10.99170017

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1549 entries, 0 to 1548
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    1538 non-null  float64
1   model                 1538 non-null  object
2   engine_power          1538 non-null  float64
3   age_in_days           1538 non-null  float64
4   km                    1538 non-null  float64
5   previous_owners       1538 non-null  float64
6   lat                   1538 non-null  float64
7   lon                   1549 non-null  object
8   price                 1549 non-null  object
9   Unnamed: 9            0 non-null     float64
10  Unnamed: 10           1 non-null     object
dtypes: float64(7), object(4)
memory usage: 133.2+ KB
```

In [5]: `df.describe()`

Out[5]:

	ID	engine_power	age_in_days	km	previous_owners	lat
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612

In [6]: `dft=df.drop(["Unnamed: 9", "Unnamed: 10"],axis=1)`

In [7]: `dff=dft.dropna()`

In [8]: `dff.isnull().sum()`

Out[8]:

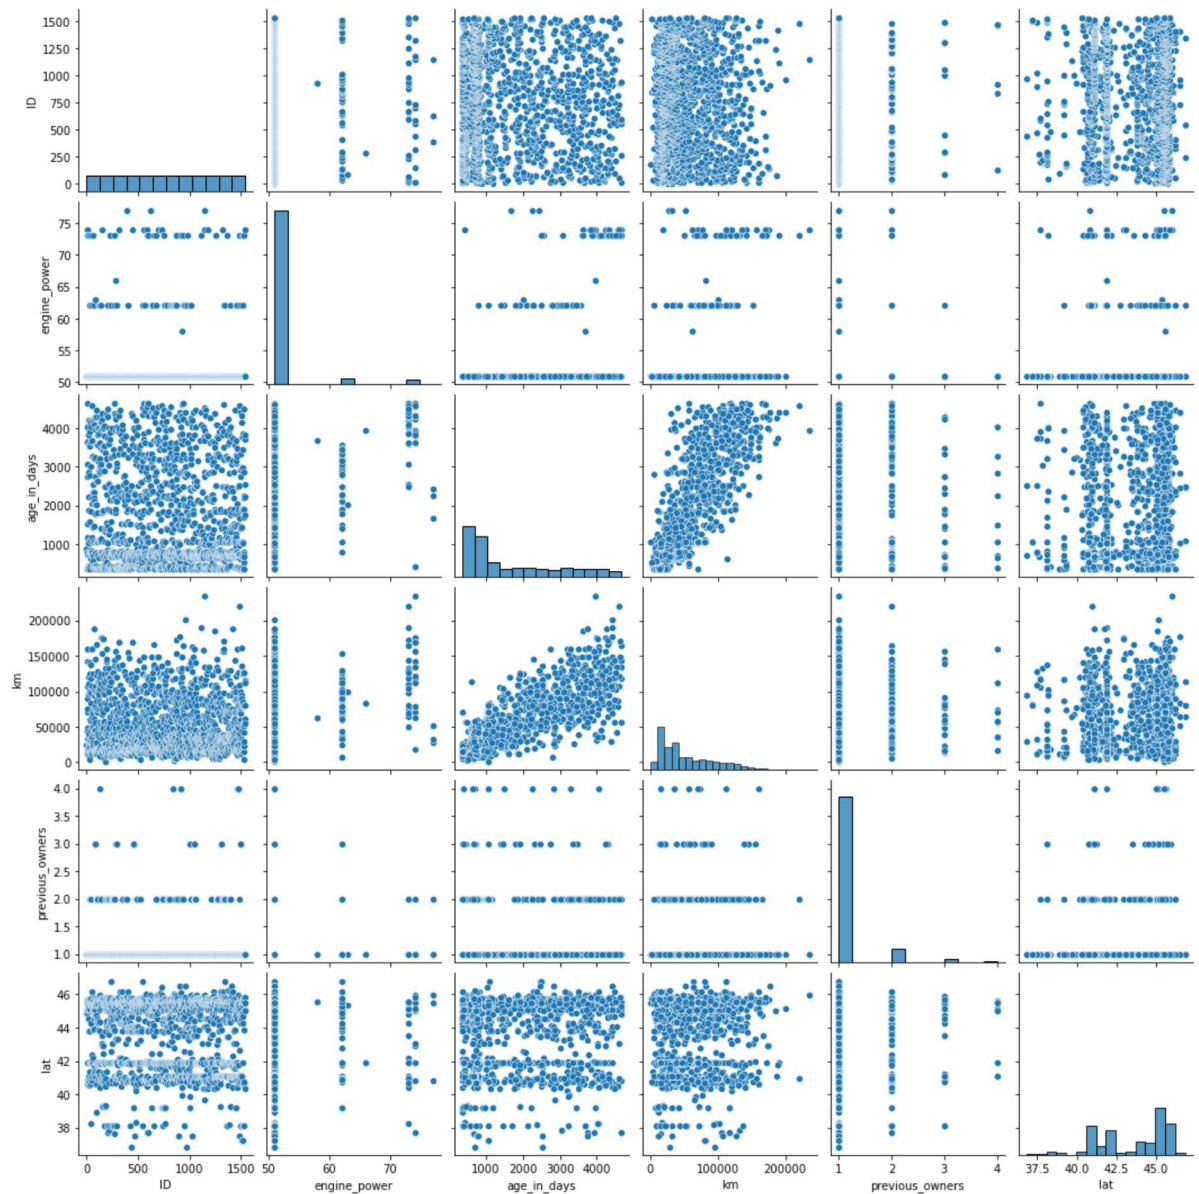
ID	0
model	0
engine_power	0
age_in_days	0
km	0
previous_owners	0
lat	0
lon	0
price	0
dtype:	int64

In [9]: `dff.columns`

Out[9]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners', 'lat', 'lon', 'price'], dtype='object')

```
In [10]: sns.pairplot(dff)
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x1ead9000ee0>
```




```
In [11]: sns.distplot(dff["km"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```


```
Out[11]: <AxesSubplot:xlabel='km', ylabel='Density'>
```



```
In [12]: df1=dfff[['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners', 'lat', 'lon', 'price']]
```

```
In [13]: sns.heatmap(df1.corr())
```

```
Out[13]: <AxesSubplot:>
```



```
In [14]: x=df1[['ID', 'engine_power', 'age_in_days', 'previous_owners',
               'lat']]
         y=df1['km']
```

```
In [15]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [16]: from sklearn.linear_model import LinearRegression
         lr=LinearRegression()
         lr.fit(x_train,y_train)
```

Out[16]: LinearRegression()

```
In [17]: print(lr.intercept_)

-3181.0780187201526
```


```
In [18]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
         coeff
```

Out[18]:

	Co-efficient
ID	5.096533
engine_power	293.020024
age_in_days	24.738277
previous_owners	5606.894794
lat	-243.470544

```
In [19]: prediction=lr.predict(x_test)
         plt.scatter(y_test,prediction)
```

Out[19]: <matplotlib.collections.PathCollection at 0x1eadb738a60>



```
In [20]: print(lr.score(x_test,y_test))

0.6947791082610598
```

```
In [21]: from sklearn.linear_model import Ridge,Lasso
```

```
In [22]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[22]: Ridge(alpha=10)
```

```
In [23]: rr.score(x_test,y_test)
```

```
Out[23]: 0.6951828056222507
```

```
In [24]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[24]: Lasso(alpha=10)
```

```
In [25]: la.score(x_test,y_test)
```

```
Out[25]: 0.6948626031682503
```

```
In [26]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[26]: ElasticNet()
```

```
In [27]: print(en.coef_)
```

```
[  5.23696939 273.84944813  24.87484581 1522.2673386 -222.90936913]
```

```
In [28]: print(en.intercept_)
```

```
1203.3292278990484
```

```
In [29]: print(en.predict(x_train))
```

```
[26696.67991866 38412.21737543 30675.50598894 ... 32820.37662363
 31448.44004584 32088.92888378]
```

```
In [30]: print(en.score(x_train,y_train))
```

```
0.6960249553313316
```

```
In [31]: from sklearn import metrics
```

```
In [32]: print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolytre Error: 16187.181219200658
```

```
In [33]: print("Mean Square Error:",metrics.mean_squared_error(y_test,prediction))
```

```
Mean Square Error: 526360230.7799673
```

```
In [34]: print("Root Mean Square Error:", np.sqrt(metrics.mean_absolute_error(y_test, pred)))
```

Root Mean Square Error: 127.2288537211613

```
In [ ]:
```