

D19

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: df=pd.read_csv(r"C:\Users\user\Downloads\22_countries.csv")
df
```

Out[3]:

	id	name	iso3	iso2	numeric_code	phone_code	capital	currency	currency_na
0	1	Afghanistan	AFG	AF	4	93	Kabul	AFN	Afghan afgh
1	2	Aland Islands	ALA	AX	248	+358-18	Mariehamn	EUR	E
2	3	Albania	ALB	AL	8	355	Tirana	ALL	Albanian
3	4	Algeria	DZA	DZ	12	213	Algiers	DZD	Algerian di
4	5	American Samoa	ASM	AS	16	+1-684	Pago Pago	USD	US Do
...	...	...	...	...	...	...	...	...	
245	243	Wallis And Futuna Islands	WLF	WF	876	681	Mata Utu	XPF	CFP fra
246	244	Western Sahara	ESH	EH	732	212	El-Aaiun	MAD	Moroco Dirh
247	245	Yemen	YEM	YE	887	967	Sanaa	YER	Yemeni
248	246	Zambia	ZMB	ZM	894	260	Lusaka	ZMW	Zamb kwa
249	247	Zimbabwe	ZWE	ZW	716	263	Harare	ZWL	Zimbat Do

250 rows × 19 columns

```
In [4]: df.head(10)
```

Out[4]:

	id	name	iso3	iso2	numeric_code	phone_code	capital	currency	currency_name
0	1	Afghanistan	AFG	AF	4	93	Kabul	AFN	Afghan afghani
1	2	Aland Islands	ALA	AX	248	+358-18	Mariehamn	EUR	Euro
2	3	Albania	ALB	AL	8	355	Tirana	ALL	Albanian lek
3	4	Algeria	DZA	DZ	12	213	Algiers	DZD	Algerian dinar
4	5	American Samoa	ASM	AS	16	+1-684	Pago Pago	USD	US Dollar
5	6	Andorra	AND	AD	20	376	Andorra la Vella	EUR	Euro
6	7	Angola	AGO	AO	24	244	Luanda	AOA	Angolan kwanza
7	8	Anguilla	AIA	AI	660	+1-264	The Valley	XCD	East Caribbean dollar
8	9	Antarctica	ATA	AQ	10	672	NaN	AAD	Antarctican dollar
9	10	Antigua And Barbuda	ATG	AG	28	+1-268	St. John's	XCD	Eastern Caribbean dollar

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250 entries, 0 to 249
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    250 non-null    int64
1   name                  250 non-null    object
2   iso3                  250 non-null    object
3   iso2                  249 non-null    object
4   numeric_code          250 non-null    int64
5   phone_code            250 non-null    object
6   capital               245 non-null    object
7   currency              250 non-null    object
8   currency_name         250 non-null    object
9   currency_symbol       250 non-null    object
10  tld                   250 non-null    object
11  native                249 non-null    object
12  region                248 non-null    object
13  subregion             247 non-null    object
14  timezones             250 non-null    object
15  latitude              250 non-null    float64
16  longitude             250 non-null    float64
17  emoji                250 non-null    object
18  emojiU               250 non-null    object
dtypes: float64(2), int64(2), object(15)
memory usage: 37.2+ KB
```

In [6]: `dff=df.dropna()`

In [7]: `dff.describe()`

Out[7]:

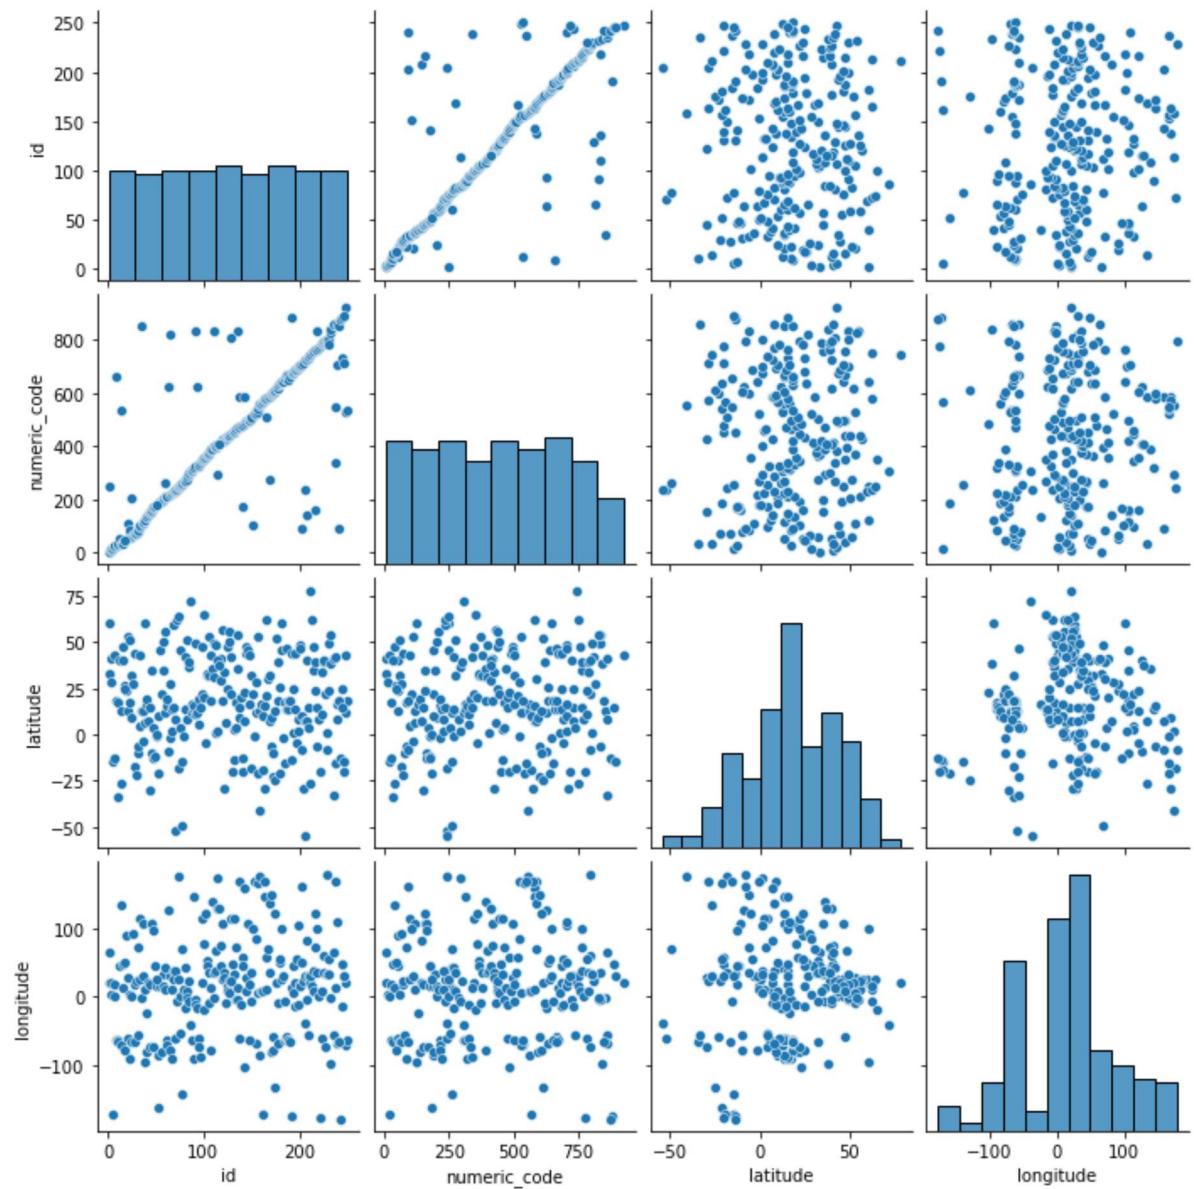
	id	numeric_code	latitude	longitude
<b>count</b>	243.000000	243.000000	243.000000	243.000000
<b>mean</b>	125.839506	437.366255	17.719476	14.241033
<b>std</b>	71.920662	254.274551	25.491128	73.423927
<b>min</b>	1.000000	4.000000	-54.500000	-176.200000
<b>25%</b>	64.500000	220.000000	1.708333	-54.000000
<b>50%</b>	126.000000	438.000000	17.000000	18.500000
<b>75%</b>	187.500000	656.500000	39.250000	49.775000
<b>max</b>	250.000000	926.000000	78.000000	178.000000

In [8]: `dff.columns`

Out[8]: Index(['id', 'name', 'iso3', 'iso2', 'numeric\_code', 'phone\_code', 'capital', 'currency', 'currency\_name', 'currency\_symbol', 'tld', 'native', 'region', 'subregion', 'timezones', 'latitude', 'longitude', 'emoji', 'emojiU'], dtype='object')

```
In [9]: sns.pairplot(dff)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x1e813ac6c40>
```




```
In [10]: sns.distplot(dff['longitude'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```


```
Out[10]: <AxesSubplot:xlabel='longitude', ylabel='Density'>
```



```
In [11]: df1=dfff[['id', 'numeric_code', 'latitude', 'longitude']]
```

```
In [12]: sns.heatmap(df1.corr())
```

```
Out[12]: <AxesSubplot:>
```



```
In [13]: x=df1[['id', 'numeric_code', 'latitude']]
         y=df1['longitude']
```

```
In [14]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [15]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[15]: LinearRegression()

```
In [16]: print(lr.intercept_)

2.2949353281480604
```


```
In [17]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[17]:

Co-efficient	
<b>id</b>	0.231902
<b>numeric_code</b>	-0.042811
<b>latitude</b>	0.094317

```
In [18]: prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[18]: <matplotlib.collections.PathCollection at 0x1e819dad490>



```
In [19]: print(lr.score(x_test,y_test))

-0.026057481797213233
```

```
In [20]: from sklearn.linear_model import Ridge,Lasso
```

```
In [21]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[21]: Ridge(alpha=10)

```
In [22]: rr.score(x_test,y_test)
```

```
Out[22]: -0.02605691270360877
```

```
In [23]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
```

```
Out[23]: Lasso(alpha=10)
```

```
In [24]: la.score(x_test,y_test)
```

```
Out[24]: -0.02592857305329921
```

```
In [25]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[25]: ElasticNet()
```

```
In [26]: print(en.coef_)
```

```
[ 0.23149464 -0.04270898  0.09334494]
```

```
In [27]: print(en.intercept_)
```

```
2.317344560640974
```

```
In [28]: print(en.predict(x_train))
```

```
[ 18.39794438  6.54194112 44.48869528 15.49183606 13.67795579
  8.49973252 17.27053298 13.296066  5.31434681 -1.68391609
 25.79699725 11.32570962  6.65092023  7.49072241 21.19243126
 19.30668588 11.16738076  5.7118749  34.47900164  1.45352994
 15.38860943 10.43404317  9.97797412 12.38885651 47.76577559
 24.15238466 19.95046058 21.25578117 14.27281983 20.43340662
  7.20067207 13.59223964  6.09073801  5.09352325 19.22577732
 46.97400677  9.61199095  1.50043289 44.97155253 23.13464965
 14.35116656 18.77071502 17.84383872 22.89694336 19.00243773
 11.65144693 10.08784171 15.94710784  1.74597981 23.69119885
 16.60382023 11.65041918 13.84663503 16.92378014 14.16875968
 10.53602128  8.82328256 14.7638074  -7.49033979  7.43306444
 19.56652039 16.90417521  0.32337025  9.43623706  6.63717342
 22.58819407 16.99390022  5.49213146  7.63575484 15.24776368
 24.04750282 15.80289061 12.31937819 20.26959998  3.28800125
 24.41529497  5.9554124  32.67267446 16.66537274 -15.05113361
 19.67326274 13.6986944  11.40558218 16.34815056  6.82057261
  7.14188931 10.68137631  6.49729941 15.65121177  4.94109079
 18.20167411 23.04269673  7.72334185 12.82590369 20.83870296
 55.89880445 34.88489105 23.64789448 19.6999564  14.85613524
  4.67235418 13.6022634  15.47757295  9.3038453  18.32028076
 17.4837074  18.27176021 19.05769034 23.74836402  5.45838647
 21.53431904 29.31801576 17.13792142  5.06712195 29.82601751
 22.90053539 27.04999505 20.9650549 -24.84032537 24.21858231
 15.11380554 16.48385262 16.50686792 38.41207368  9.40295545
 13.73756713  3.75812681  7.82043048 15.46317871 16.93010252
  9.5506718  18.02476412  4.51458991 21.61555853  6.71427741
  9.93109736 22.55084367 10.80393757  6.4359555  13.03306769
  8.47775084 16.84076271 16.97425306  6.81929303 19.9159912
 39.06773225 -16.27029825  9.7635587 -22.31507818 11.99311049
 -3.15487526  5.66436758 15.63624626  5.34447386 18.46014338
 21.78518714  8.56032524 21.57618563  9.15148854 21.09257918
 13.82335242 12.33293901 15.72712908 15.68382471 18.61661672
 32.2835366  11.90350924  7.34260869 10.88052768 14.75980263]
```

```
In [29]: print(en.score(x_train,y_train))
```

```
0.02038555333192482
```

```
In [30]: from sklearn import metrics
```

```
In [31]: print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolytre Error: 46.24951625068402
```

```
In [32]: print("Mean Square Error:",metrics.mean_squared_error(y_test,prediction))
```

```
Mean Square Error: 4282.361608652522
```

```
In [33]: print("Root Mean Square Error:",np.sqrt(metrics.mean_absolute_error(y_test,pre
```

```
Root Mean Square Error: 6.800699688317668
```



```
In [34]: import pickle
```

```
In [35]: f4="prediction"  
pickle.dump(lr,open(f4,'wb'))
```

```
In [ ]:
```